

Algorand Protocol Description

Argimiro, CoinFabrik

March 11, 2023

Contents

1	Introduction	1
2	Player accounts	2
3	Ledger definition	2
4	Node as a State Machine	2
5	Main algorithm	3
5.1	Block Proposal	4
5.2	Soft Vote	5
5.3	Proposal handler	6
5.4	Vote handler	7
5.5	Bundle handler	7
5.6	Recovery Attempt	8
5.7	Fast Recovery Attempt	8
6	Subroutines	9

1 Introduction

We aim to describe how the Algorand protocol works. This document is meant to reflect how the Algorand mainnet is behaving today. We relied on Algorand's documents ([GHM⁺17], [Mic16] and Algorand's official specifications) but consulted the node's code and probed the network when information was unclear or unavailable.

Algorand is a proof-of-stake blockchain cryptocurrency protocol using a decentralized Byzantine Agreement protocol that leverages pure proof of stake for consensus. The protocol maintains a ledger that is modified via consensus. In particular, this ledger encodes how many ALGOs (the native token) holds each account. Blocks encode the status of the ledger. Starting on the genesis block (round 0) that encodes the first state of the network, on each round the participation nodes vote on what will the next block be. Their voting power is proportional to the stake (in ALGOs) held by the accounts associated to the node.

At the genesis of the Algorand blockchain, 10Bn ALGO was minted. As of September 2022, circulating supply is approximately 6.9b ALGO, distributed through different forms of ecosystem support and community incentives. The remaining Algos are held by the Foundation in secure wallets assigned to Community and Governance Rewards (%53), Ecosystem Support (%36), and Foundation endowment (%11).

A network is formed with two kind of nodes: relay and participation nodes. Participation nodes can connect only to relay nodes. They listen to one or more participation nodes and they may send a message to a relay node. Relay nodes simply listen to participation nodes connected to them and relay the messages they receive.

Explain how
the network
is formed

In the following document we present a detailed pseudocode interpretation of the program run by an Algorand node managing at least one online (i.e., participating) account. We attempt to explain the different states the node can be in, as well as all kinds of events that cause transitions and network output. This work aims to provide a basis for an Algorand blockchain simulator currently in development.

2 Player accounts

A player is a unique participant of the protocol. We will use the terms player and account indistinctly throughout this document. A player is uniquely identified by a 256-bit string I called an address. Furthermore, they possess a certain balance (an unsigned 64 bit integer), and a pair of participation keys set up according to the key scheme. An account may be online, offline or nonparticipating. An account is offline when it has not sent an online transaction, that creates a participation key and a set of ephemeral keys to be used at each round and destroyed at the end. Also, an account has a set of VRF keys, computed at each round. TODO: keep describing player. Describe register process, online / offline / nonpart attributions

3 Ledger definition

A ledger is a sequence of states which comprise the common information established by some instantiation of the Algorand protocol. Given a Ledger instance, its immutable characteristics are (...) A ledger state then is defined by the following parameters:

A Ledger entry e (also known as block) is a data structure that models a ledger's state transition.

4 Node as a State Machine

We model each player as a finite state machine. In order to do so, we will first define and explain all variables involved in a given state a player is able to be in or transition to. A node, the network entity running the protocol, will then be "playing for" one or more online accounts that it manages.

Let round and period number, r and p , be unsigned 64 bit integers. Let the step number, s , be an unsigned 8 bit integer. For convenience and readability, the following aliases for step number are defined:

- $s = 0 \equiv \textit{proposal}$
- $s = 1 \equiv \textit{soft}$
- $s = 2 \equiv \textit{cert}$
- $s \in [3, 252] \equiv \textit{next}_{s-3}$
- $s = 253 \equiv \textit{late}$
- $s = 254 \equiv \textit{redo}$
- $s = 255 \equiv \textit{down}$

A proposal-value is a tuple $v = (I, p, \text{Hash}(e), \text{Hash}(\text{Encoding}(e)))$ where I is an address (the "original proposer"), p is a period (the "original period"), Hash is some cryptographic hash function (usually SHA512/256), Encoding is performed in msgpack, and e is a given ledger entry (as described in the previous section). The special proposal where all fields are the zero-string is called the bottom proposal \perp .

Consider now tuples of the form (I, r, p, s, v) , where v is a value. We will call this kind of construct a vote, noting it $\text{Vote}(r, p, s, v)$, read as a vote for value ' v ' at round r , period p and step s .

We now consider ledger entries of the form $(...)$. A tuple of the form $()$ is called a proposal or proposal payload, noted $\text{Proposal}(v)$.

Given all of the above data types and constructs, a node can be modelled as a finite state machine. A state is given by a 3-tuple of integers, (r, p, s) , two sets P and V of objects which we will call observed *proposal* and *vote* values, a special value v_{prime} which we call the pinned value (see recovery algorithms), a Ledger L and a balance table BT , a table of public addresses for all accounts holding any stake in the network, and their balances. A Ledger is an ordered chain of entries, and fundamentally it represents state changes in the balance table. The round number can be thought of as an index into the ledger, where the event of observing a new round cannot be decoupled from the event of adding an entry (block) to the ledger.

5 Main algorithm

Algorithm 1 Main node algorithm

```

1: function EventHandler(Event  $ev$ )
2:   if  $ev$  is TimeoutEvent then
3:      $time \leftarrow ev.time$ 
4:     if  $time = 0$  then
5:       BlockProposal()
6:     else if  $time = \text{FilterTimeout}(p)$  then
7:       SoftVote()
8:     else if  $time = \max\{4\lambda, A\} \vee time = \max\{4\lambda, A\} + 2^{st-3}\lambda + r, st \in [4, 252], r \in$ 
9:        $[0, 2^{st-3}\lambda]$  then
10:      Recovery()
11:    else if  $\exists k \in \mathbb{Z}, r \sim U(0, \lambda_f), time = k\lambda_f + r$  then
12:      FastRecovery()
13:    else
14:       $msg \leftarrow ev.msg$ 
15:      if  $msg$  is Proposal  $p$  then
16:        HandleProposal( $p$ )
17:      else if  $msg$  is Vote  $v$  then
18:        HandleVote( $v$ )
19:      else if  $msg$  is Bundle  $b$  then
20:        HandleBundle( $b$ )

```

Algorand nodes operate based on events, which are the only means through which the node state machine can transition and produce outputs. Whenever an event is received, the node will produce a state change and will typically generate an output that will be broadcasted or relayed to other nodes on the network.

There are two types of events: Timeout events and Message events. Timeout events are produced when the node's internal clock reaches a certain time since the start of the current period, while Message events are generated in response to messages sent by other network actors.

The main algorithm describes how the event handling mechanism works. If the event is a timeout and a new period has just started, the node will run the block proposal step. After a certain amount of time, the node will attempt to soft vote, selecting the best received proposal value based on the lowest proposer hash criteria and broadcasting its votes to the network if selected as part of the soft voting committee.

If the event is a message reception, the node will handle it according to its context. For example, a soft vote outside of the soft voting step could produce a next vote, but it could not produce a certification vote without a period change and another round of voting.

After a certain amount of time, the node will start the recovery steps and next vote a value or an empty block if it observes the required threshold votes. Finally, after another specified amount of time, the block will attempt a FastRecovery.

More details on the specifics of each handler can be found in their respective subsections.

5.1 Block Proposal

Algorithm 2 Block proposal

```

1: function BlockProposal
2:   ResynchronizationAttempt()
3:   for  $a \in A$  do
4:      $\langle \text{sorthash}, \pi, j \rangle \leftarrow \text{Sortition}(a_{sk}, \text{getSortitionSeed}(ctx, r), t, \text{proposal},$ 
                                      $\text{getSortition}_w(ctx, r, a_{pk}), \text{getSortition}_W(ctx, r))$ 
5:     if  $j > 0$  then
6:       if  $p = 0 \vee \exists s' > 2/Bundle(r, p - 1, s', \perp) \subset V$  then
7:          $e \leftarrow \text{AssembleBlock}()$ 
8:          $v \leftarrow \text{Proposal}_{value}(e)$ 
9:          $\text{Broadcast}(\text{Vote}(I, r, p, \text{proposal}, v))$ 
10:         $\text{Broadcast}(e)$ 
11:       else
12:          $v \leftarrow v_0 | \exists s, Bundle(r, p - 1, s, v_0) \subset V$ 
13:          $\text{Broadcast}(\text{Vote}(I, r, p, \text{proposal}, v))$ 
14:         if  $\text{Proposal}(v) \in P$  then
15:            $\text{Broadcast}(\text{Proposal}(v))$ 

```

Description:

The node manages multiple accounts and executes a "play" for each account, as each account represents a distinct player in the system. The node runs a sortition process for the proposal step of each account. If the account is selected to propose (i.e., if the account has more than 0 votes for the specific context), the node proceeds with the proposal step and period.

If the check on line 5 passes (i.e., if it is the first period or an empty block has been observed during the recovery step of the previous period), the node assembles a new block (e), calculates the proposal value (v), and broadcasts a vote for v on step 0 along with the full block (e).

If the check on line 5 fails (i.e., it is not the first period and no empty block was observed during the recovery step of the previous period), the period change must have been triggered by observing a bundle for a specific value v_0 . The node broadcasts a vote for this value.

If the actual proposal (i.e., the full ledger entry) is available, the node then broadcasts it in a separate message.

5.2 Soft Vote**Algorithm 3** *Soft Vote*

```

1: function SoftVote
2:    $lowestObservedHash \leftarrow \infty$ 
3:    $v \leftarrow \perp$ 
4:   for  $vote \in V, vote.step = proposal$  do
5:      $priorityHash \leftarrow \min_{i \in [0, j)} Hash(vote.sorthash || i)$ 
6:     if  $priorityHash < lowestObservedHash$  then
7:        $lowestObservedHash \leftarrow priorityHash$ 
8:        $v \leftarrow vote.v$ 
9:   for  $a \in Accounts$  do
10:     $\langle j, hash, \pi \rangle \leftarrow sortition(a, soft)$ 
11:    if  $j > 0 \wedge lowestObservedHash < \infty$  then
12:       $Broadcast(Vote(r, p, soft, v))$ 
13:      if  $Proposal(v) \in P$  then
14:         $Broadcast(Proposal(v))$ 

```

Description:

The soft vote routine (also known as "filtering") is triggered after a SV timeout since the beginning of a period. Of all received proposal votes, the node finds the one that minimizes the priority function. The priority function is computed, for each voting unit of each proposer, as the minimum $Hash(sorthash \text{ --- } i)$, where sorthash is the proposer's sortition output hash, i is an integer in between 1 (inclusive) and the j value output by the proposer's sortition run (exclusive). If no valid votes were observed by the node, there is no output and the algorithm ends here. If a value with minimum priority is found, every account managed by the node runs sortition to be part of the soft vote committee. If selected, the vote is broadcast. If the ledger entry associated to that value has also been observed, it is broadcast separately but immediately after.

5.3 Proposal handler

Algorithm 4 HandleProposal

```

1: function HandleProposal(Proposal e)
2:    $v \leftarrow \text{Proposal}_{\text{value}}(e)$ 
3:
4:   if  $\sigma(S, r + 1, 0) = v$  then
5:     Relay(e)
6:
7:   if  $\text{isValid}(e) \wedge e \notin P \wedge v \in \{\sigma(S, r, p), v, \mu(S, r, p)\}$  then
8:     Relay(e)
9:      $P \leftarrow P \cup \{e\}$ 
10:
11:    if  $\text{IsCommittable}(v) \wedge s \leq \text{cert}$  then
12:      for  $a \in A$  do
13:         $\langle \text{sorthash}, \pi, j \rangle \leftarrow \text{Sortition}(a_{sk}, \text{getSortitionSeed}(\text{ctx}, r), t, \text{cert},$ 
14:                                      $\text{getSortition}_w(\text{ctx}, r, a_{pk}), \text{getSortition}_W(\text{ctx}, r))$ 
15:        if  $j > 0$  then
16:          Broadcast(Vote(a.I, r, p, cert, v))

```

Description:

The proposal handler is activated when a message event for a full proposal (aka. a ledger entry, *e*) is received by the node. There's two parts to the algorithm. First, a series of checkups are performed (...) Afterwards, and assuming the proposal is valid and not to be ignored, a check for commitability of the proposal value is performed. If the value is commitable (that is, there is a soft bundle for it - more on that on `IsCommittable()` -), and the current step number is `cert` (2) or less, we run sortition for the account for a membership of the certification committee. If selected, a certification vote is broadcast.

5.4 Vote handler

Algorithm 5 HandleVote

```

1: function HandleVote(Vote vt)
2:   if  $\neg \text{VerifyVote}(vt) \vee (s = 0 \wedge vt \in V) \vee (s = 0 \wedge \text{IsEquivocation}(vt)) \vee (s > 0 \wedge$ 
      $\text{IsSecondEquiv}(vt)) \vee vt.r \notin [r, r+1] \vee (vt.r = r+1 \wedge (vt.p > 0 \vee vt.s \in (next_0, late))) \vee (vt.r =$ 
      $r \wedge (vt.p \notin [p-1, p+1] \vee vt.p = p+1 \vee (vt.p = p \wedge vt.s \in (next_0, late) \wedge vt.s \notin [s-1, s+1]) \vee (vt.p =$ 
      $p-1 \wedge vt.s \in (next_0, late) \wedge vt.s \notin [hs-1, hs+1]))$  then
3:     return //the vote is ignored
4:   //if we got here, the vote is to be relayed and observed
5:   Relay(vt)
6:    $V \leftarrow V \cup vt$ 
7:   if  $vt.s = proposal$  then //reproposal payload
8:     Broadcast(Proposal(vt.v))
9:   else if  $vt.s = soft \wedge \text{IsCommitable}(vt.v) \wedge s \leq cert$  then
10:    for Account a  $\in A$  do
11:       $\langle sorthash, \pi, j \rangle \leftarrow \text{Sortition}(a_{sk}, \text{getSortitionSeed}(ctx, r), t, cert,$ 
         $\text{getSortition}_w(ctx, r, a_{pk}), \text{getSortition}_W(ctx, r))$ 
12:      if  $j > 0$  then
13:        Broadcast(Vote(a.I, r, p, cert, vt.v))
14:    else if  $vt.s = cert \wedge \text{Bundle}(r, p, cert, vt.v) \subset V$  then
15:      Commit(Proposal(vt.v))
16:      ObserveNewRound()
17:    else if  $vt.s \in [next_0, next_{250}] \wedge \exists v | \text{Bundle}(r, p, vt.s, v) \subset V$  then
```

Description:

The vote handler is called whenever a node receives a vote of any kind.

5.5 Bundle handler

Algorithm 6 HandleBundle

```

1: function HandleBundle(Bundle b)
2:   if  $\text{VerifyBundle}(b) \wedge b.r = r \wedge b.p \geq p-1$  then
3:     for vote  $\in b$  do
4:       HandleVote(vote)
```

Description:

This handler is called when a message containing a bundle is received. If the bundle is verified (which implies verifying the validity of every vote separately), the round of all votes in it matches the node's current round and its period is at least the previous one ($p-1$), then each vote in the bundle is processed independantly. Functionally, receiving a bundle is the same as receiving a set of votes sequentially.

5.6 Recovery Attempt

Algorithm 7 Recovery

```

1: function Recovery
2:    $s \leftarrow next_s$ 
3:   ResynchronizationAttempt()
4:   for  $Account a \in A$  do
5:      $\langle sorthash, \pi, j \rangle \leftarrow \text{Sortition}(a_{sk}, \text{getSortitionSeed}(ctx, r), t, next_s,$ 
6:                                      $\text{getSortition}_w(ctx, r, a_{pk}), \text{getSortition}_W(ctx, r))$ 
7:     if  $j > 0$  then
8:       if  $\exists v | IsCommitable(v)$  then
9:         GOSSIP(Vote( $I, r, p, next_s, v$ ))
10:      else if  $\nexists s_0 > cert | Bundle(r, p-1, s_0, \perp) \subseteq V \wedge$ 
11:               $\exists s_1 > cert | Bundle(r, p-1, s_1, \bar{v}) \subseteq V$  then
12:        GOSSIP(Vote( $I, r, p, next_s, \bar{v}$ ))
13:      else
14:        GOSSIP(Vote( $I, r, p, next_s, \perp$ ))

```

Description:

5.7 Fast Recovery Attempt

Algorithm 8 FastRecovery

```

1: function FastRecovery
2:   ResynchronizationAttempt()
3:   for  $Account a \in A$  do
4:     if  $IsCommitable(v)$  then
5:        $\langle sorthash, \pi, j \rangle \leftarrow \text{Sortition}(a_{sk}, \text{getSortitionSeed}(ctx, r), t, 'late',$ 
6:                                      $\text{getSortition}_w(ctx, r, a_{pk}), \text{getSortition}_W(ctx, r))$ 
7:       if  $j > 0$  then
8:         GOSSIP(Vote( $I, r, p, late, v$ ))
9:       else if  $\nexists s_0 > cert | Bundle(r, p-1, s_0, \perp) \subseteq V \wedge$ 
10:               $\exists s_1 > cert | Bundle(r, p-1, s_1, \bar{v}) \subseteq V$  then
11:          $\langle sorthash, \pi, j \rangle \leftarrow \text{Sortition}(a_{sk}, \text{getSortitionSeed}(ctx, r), t, 'redo',$ 
12:                                      $\text{getSortition}_w(ctx, r, a_{pk}), \text{getSortition}_W(ctx, r))$ 
13:         if  $j > 0$  then
14:           GOSSIP()
15:       else

```

Description:

The fast recovery algorithm is executed periodically every integer multiple of λ_F seconds. In it, nodes of the network make an attempt to reestablish normal functioning of the network by broadcasting, in the hopes of observing, any kind of consensus. Firstly, and same as in the previous mentioned "next" recovery stages, a resynchronization attempt is made (broadcasting freshest bundle or value if there is any, see *ResynchronizationAttempt*() in the next section).

6 Subroutines

Algorithm 9 IsCommitable

```

1: function IsCommitable( $Proposal_{value} v$ )
2:   return  $Proposal(v) \in P \wedge Bundle(r, p, soft, v) \subset V$ 

```

Arguments:

- v = a proposal value (as we know from a previous section, computed as a function of a ledger entry in consideration).

Description:

A boolean helper subroutine that returns *True* if, for the input value v , a ledger entry $e = Proposal(v)$ has been observed by the player and is available in the current state, and a soft bundle of votes b has been observed, where for every vote $vt \in b$ either $vt.v = v$ or vt is an equivocation vote (in which case it counts as a "wildcard" towards any value; in particular v in this case). Otherwise, the subroutine returns *False*.

Algorithm 10 ResynchronizationAttempt

```

1: function ResynchronizationAttempt
2:    $Val = \perp$ 
3:   if  $\exists v | Bundle(r, p, soft, v) \subset V$  then
4:      $Broadcast(Bundle(r, p, soft, v))$ 
5:      $val = v$ 
6:   else if  $\exists s_0 > cert | Bundle(r, p - 1, s_0, \perp) \subset V$  then
7:      $Broadcast(Bundle(r, p, s_0, \perp))$ 
8:   else if  $\exists s_0 > cert, v \neq \perp | Bundle(r, p - 1, s_0, v) \subset V$  then
9:      $Broadcast(Bundle(r, p, s_0, v))$ 
10:     $val = v$ 
11:   if  $val \neq \perp$  and  $Proposal(v) \in P$  then
12:      $Broadcast(Proposal(v))$ 

```

Description:

Algorithm 11 ObserveNewRound

```

1: function ObserveNewRound
2:    $\bar{s} \leftarrow s$ 
3:    $\bar{v} \leftarrow \perp$ 
4:    $s \leftarrow 0$ 
5:    $r \leftarrow r + 1$ 
6:    $p \leftarrow proposal$ 
7:    $GarbageCollect()$ 

```

Description:

Algorithm 12 ObserveNewRound

```
1: function ObserveNewPeriod
2:    $\bar{s} \leftarrow s$ 
3:    $s \leftarrow 0$ 
4:    $p \leftarrow p + 1$ 
5:   GarbageCollect()
```

Description:

Algorithm 13 IsEquivocation

```
1: function IsEquivocation( $Vote_{vt}$ )
   return  $\exists v' \neq vt.v \mid Vote(vt.I, vt.r, vt.p, vt.s, vt.v) \in V \wedge Vote(vt.I, vt.r, vt.p, vt.s, v') \in V$ 
```

Description:

Algorithm 14 ComputeSeed

```
1: function ComputeSeed
```

Arguments:

- x

Description:

This subroutine computes the required sortition seed for the given round number, which goes in the proposed block's metadata. If the block is empty, the seed is a hash of the previous block's seed. The $get_{SK_a}(ctx, r)$ helper function gets the relevant account's secret ephemeral keys (according to the signing scheme described in specs, the keys 160 rounds prior to r). This roughly corresponds to the secret key from a round b time before block $r - 1 - (r \bmod R)$, where R is the sortition seed's renewal rate, r is the current round's number, and b is the upper bound for the maximum ammount of time that the network might be compromised.

Returns:

- the computed seed for the given block, ledger context and round

Algorithm 15 getSortitionSeed

```
1: function getSortitionSeed( $ctx, r, a_{pk}$ )
   return  $ctx.block[r - 1 - (r \bmod R)].seed$ 
```

Arguments:

- ctx = a helper structure to retrieve ledger context information (e.g. the last confirmed block)
- $round$ = current round number
- a_{pk} = the account's public key for the look up table

Description:

This helper function gets the relevant sortition seed for the current round r , according to the seed lookback parameter R . Conceptually, it corresponds with the seed computed R rounds prior to r , refreshed every R rounds.

Returns:

- a sortition seed to be used in the round r

Algorithm 16 getSortitionWeight

```

1: function getSortitionw(ctx, round, apk)
    return ctx.balanceTable[r - (R + SL)] [apk]

```

Arguments:

- ctx = a helper structure to retrieve ledger context information (e.g. the last confirmed block)
- r = current round number
- a_{pk} = the account's public key for the look up table

Description:

This helper function retrieves the stake $R + SL$ rounds prior to r , for an account with public key a_{pk}

Returns:

- the relevant account's stake

Algorithm 17 getSortitionTotalStake

```

1: function getSortitionW(ctx, r)
    return  $\sum_{a_{pk} \in ctx.balanceTable[r-(R+SL)]} balanceTable[r-(R+SL)][a_{pk}]$ 

```

Arguments:

- ctx = a helper structure to retrieve ledger context information (e.g. the last confirmed block)
- r = current round number

Description:

This helper function returns the sum of all stake for $R + SL$ rounds prior to r .

Returns:

- the total stake at play in the relevant round (according to lookback parameters)

Algorithm 18 Sortition

```

1: function Sortition(sk, seed,  $\tau$ , role, w, W)
2:    $\langle hash, \pi \rangle \leftarrow \text{VRF}(seed || role)$ 
3:    $p \leftarrow \frac{t}{W}$ 
4:    $j \leftarrow 0$ 
5:   while  $\frac{hash}{2^{hashlen}} \notin [\sum_{k=0}^j \mathcal{B}(k; w, p), \sum_{k=0}^{j+1} \mathcal{B}(k; w, p)]$  do
6:      $j++$ 
    return  $\langle hash, \pi, j \rangle$ 

```

Arguments:

- sk = a user's secret key (an ephemeral key for the given round, according to key specs)
- $seed$ = the sortition seed to be used
- τ = the expected committee size for the given role

- $role$ = a flag specifying the role for the sortition to take place (e.g. block proposer)
- w = the user's weight (i.e., its relevant stake)
- W = the total relevant stake for the given round

Description:

The Sortition procedure is one of the most important subroutines in the main algorithm, as it is used in multiple stages and contexts. Generally, it manages to pseudo-randomly but verifiably (through the use of a Verifiable Random Function) select a user with probability proportional to their weight (i.e., stake) by returning a j parameter, which indicates the number of times that specific user was chosen. Algorithmically, every monetary unit the user has is considered a "sub-user", and then each one of them is selected with probability $p = \frac{\tau}{W}$, where τ is the expected amount of users to be selected for the given role. The semi-open interval $[0, 1)$ is then split into consecutive intervals using an accumulated binomial distribution, and wherever the fraction $\frac{hash}{2^{hashlen}}$ belongs to the interval, that's the ammount of chosen sub-users for the subroutine caller.

Returns:

- an integer j that will be positive (larger than 0) if the user has been selected, and its size corresponds to the amount of sub-users for a given committee member

Algorithm 19 Commit

```

1: function Commit(LedgerEntry  $e$ )
2:    $L \leftarrow L || e$ 
3:   GarbageCollect()
4:    $r \leftarrow r + 1$ 
5:   UpdateBalanceTable()

```

Arguments:

- *LedgerEntry* = a ledger entry deemed to be committable by the consensus mechanism. That is, e is a valid entry for L , and a valid certification bundle has been observed for its corresponding value (the unique value v for which $Proposal_{value}(e) = v$)

Description:

This subroutine commits an entry e to the ledger. The entry passed as argument is appended to the ledger L , the current state is garbage collected, a new round is observed and the balance table and general account records are updated with whatever transaction information was present in e .

Algorithm 20 GarbageCollect

```

1: function GarbageCollect
2:    $P \leftarrow P - \{\forall e \in P, e.r < r\} - \{\forall e' \in P, e'.r = r \wedge e'.p + 1 < p\}$ 
3:    $V \leftarrow V - \{\forall vt \in V, vt.r < r\} - \{\forall vt' \in V, vt'.r = r \wedge vt'.p + 1 < p\}$ 

```

Description:

Algorithm 21 VerifySortition

```
1: function VerifySortition( $pk, seed, \tau, role, w, W$ )
2:   if  $\neg \text{VerifyVRF}_{pk}(hash, \pi, seed || role)$  then return 0
3:    $p \leftarrow \frac{t}{W}$ 
4:    $j \leftarrow 0$ 
5:   while  $\frac{hash}{2^{hashlen}} \notin [\sum_{k=0}^j B(k; w, p), \sum_{k=0}^{j+1} B(k; w, p))$  do
6:      $j++$ 
   return  $j$ 
```

Arguments:

- pk = a user's public key (their address)
- $seed$ = the sortition seed to be used
- τ = the expected committee size for the given role
- $role$ = a flag specifying the role for the sortition to take place (e.g. block proposer)
- w = the user's weight (i.e., its relevant stake)
- W = the total relevant stake for the given round

Description:

The sortition verification procedure takes Sortition's output and utilizes VRF properties to verify the validity of said output. Once the check is passed, it repeats Sortition's sub-user selection procedure, and outputs the amount of times the user was chosen according to their respective observed stake.

Returns:

- an integer j that will be positive (larger than 0) if the user has been selected, and its size corresponds to the amount of sub-users for a given committee member

Appendix: Notation

- R = sortition seed renewal rate (in number of rounds). Set to 2 in the specs. as of December 2022
- τ_{step} = expected number of members on a regular step committee
- τ_{final} = expected number of members on a final consensus achieving committee
- T_{step} = fraction of expected members for a voting committee on a given step
- T_{final} = fraction of expected members for a voting committee on the final step (step = 2). If observed, achieves final consensus on the given block hash
- $\lambda_{proposal}$ = time interval for the node to accept block proposals, after which it chooses the observed block with the highest priority (lowest hash)
- λ_{block} = waiting time for the full block to be received once decided by vote. If no full block is received, the node falls back to the empty block.
- SL = the account balances lookback interval, in number of rounds (integer). Set to 320 in specs.
- Public keys a_{pk} for every account linked to the node (doubles as their respective addresses)

- Public participation keys $a_{p-partkey}$ for all accounts registered as online (private participation keys are kept securely by respective users and should not be directly accesible by the node)
- Two levels of ephemeral keys, signed according to the signing scheme described in specs. (where first level keys are signed with the participation keys, and second level (aka. leaf) ephemeral keys are signed with first level ephemeral keys). These keys are used for actual participation in the consensus (abstracted as a_{sk} in the following code) They are ephemeral because they live for a single round, after which they are deleted
- A balance table *BalanceTable* for all accounts linked to the node, for the round $SL + R$ rounds before the one currently running

References

- [GHM⁺17] Yossi Gilad, Rotem Hemo, Silvio Micali, Georgios Vlachos, and Nickolai Zeldovich. Algorand: Scaling byzantine agreements for cryptocurrencies. In *Proceedings of the 26th Symposium on Operating Systems Principles, Shanghai, China, October 28-31, 2017*, pages 51–68. ACM, 2017.
- [Mic16] Silvio Micali. ALGORAND: the efficient and democratic ledger. *CoRR*, abs/1607.01341, 2016.